
Inflation versus Labor-Market Attention in FOMC Statements and Minutes

A descriptive NLP report on the Federal Reserve dual mandate, 2000–2024

Kamil Stos
ENSAE Paris
kamil.stos@ensae.fr

Abstract

This report studies how FOMC statements and minutes differ in their treatment of the two sides of the Federal Reserve’s mandate: inflation and the labor market. The corpus contains 199 meetings from 2000 to 2024, with one statement and one set of minutes for each meeting. I start from a transparent dictionary score: inflation language minus labor-market language, normalized by document length. In this sample, statements are shorter, more concentrated, and more inflation-dominant than minutes. Minutes are longer and closer to balance, with more space for labor-market conditions, uncertainty, and financial stress. I then move to sentences. A balanced sample of sentences is labelled with a local language model and used as a weakly annotated training set. Simple sentence classifiers based on TF–IDF and local embeddings are trained to distinguish inflation, labor, financial-conditions, policy-decision, and other sentences. The embedding classifier agrees more closely with the weak labels than the dictionary rule and gives a useful second view of the corpus. The results should be read descriptively: the dictionary remains the main interpretable measure, while the sentence model shows where exact keyword matching misses broader financial-conditions language.

Introduction

The Federal Reserve has a dual mandate: price stability and maximum employment. After each FOMC meeting, the Committee publishes a short statement. A few weeks later, it publishes longer minutes. Both documents refer to the same meeting, but they play different communication roles. The statement is the immediate public message; the minutes give a more detailed account of the discussion.

This report asks a simple question:

Do statements and minutes give the same relative attention to inflation and the labor market? Do they treat both subjects the same way?

We look at the topics that appear in the documents, the space given to each side of the mandate, and the way this changes across document types and periods.

The difference is interesting because the Fed’s legal mandate is dual, but price stability has a privileged place as the initial original mandate. Inflation is directly tied to credibility, expectations, and the policy instruments.

This report takes the macro-topic angle: inflation versus the labor market. The companion report by Mathieu takes the symmetric policy-tone angle: hawkish versus dovish communication. The two views complement each other. Inflation language often supports a hawkish signal, but the mapping

is not exact. Labor-market language can describe wage pressure, demand, or resilience while the overall policy stance remains restrictive.

The analysis proceeds in three steps. First, I compare statements and minutes with the dictionary score across document types, time, policy episodes, and rate decisions. This gives the main observation that statements are more inflation-heavy than minutes. Second, I move to sentences and try to understand closer what the dictionary is actually counting. Third, I use LLM-assisted weak labels to train simple sentence classifiers with TF-IDF and embedding features. This last step is more of a check on whether broader topic categories, especially financial conditions, can be recovered from sentence text. These different methods are rather used as a sequence of checks on the same object, to try to catch quirks or different points of view.

Related Work

There is an extensive body of literature on FOMC communication and especially in relation to financial markets, fixed income instruments and especially Overnight Indexed Swap (OIS) rates since they are directly translating the market’s expectation of the FED’s decision. Tadde (Tadde, 2022) builds sentiment indices for statements and minutes and relates minutes sentiment to financial-market reactions. However, in this report, we focus solely on the language to stay in the scope of the project.

There is still a large body of literature that aims at using text methods to study FOMC communication directly. Acosta and Meade (Acosta and Meade, 2015) study post-meeting statements with semantic similarity methods and document how statements evolved over time. Hansen, McMahon, and Prat (Hansen et al., 2018) use computational linguistics on internal FOMC transcripts to study deliberation and transparency inside the Committee. Their data are richer than mine because transcripts contain individual interventions. This report uses only public statements and minutes, so it measures the communication that is released to the public, not the full internal discussion.

The closest paper for the dual-mandate motivation is Shapiro and Wilson (Shapiro and Wilson, 2022). They estimate central-bank objectives from the language of internal meetings, including the relative weight placed on inflation and real activity. The goal of my report is rather to be descriptive and to get a feel for how both subjects are treated linguistically, not evaluate their economical relevance.

Doh, Kim, and Yang (Doh et al., 2021) show that qualitative descriptions in FOMC statements can matter for bond prices. Doh, Song, and Yang (Doh et al., forthcoming) use alternative FOMC statements and language models to measure hawkish and dovish policy stance. Those papers are closer to the companion report on hawkish versus dovish communication.

Data

The corpus comes from the official Federal Reserve website and is retrieved and parsed automatically. Each row corresponds to one meeting and contains the meeting date, URLs, statement text, and minutes text.

The sample contains 199 meetings from February 2000 to December 2024. Every meeting has both a statement and minutes. Rate decisions are inferred from statement text with simple regular expressions and coded as *hike*, *cut*, or *hold*. The sample has 59 hikes, 30 cuts, and 110 holds.

Table 1: Corpus size and document length.

Document type	Documents	Mean words	Median words	Std. words
Statements	199	526.1	536.0	217.4
Minutes	199	7283.2	7536.0	2441.3




Figure 1: Document length by meeting. Minutes are much longer than statements, so dictionary scores are normalized by document length.

Minutes contain about fourteen times more words than statements on average, so the main quantities below are reported as mentions per 1,000 words.

We almost don't preprocess the text, it is just lowercased, whitespace is normalized, and then words are counted with a regular expression. Stopwords are kept for the dictionary measure because several terms are multi-word expressions, such as *price stability*, *labor market*, and *inflation expectations*.

Dictionary balance

I define two short dictionaries. The inflation dictionary contains terms such as *inflation*, *inflationary*, *price stability*, *price pressures*, *inflation expectations*, *core inflation*, *PCE inflation*, *consumer prices*, *energy prices*, *food prices*, *cost pressures*, *supply constraints*, and *disinflation*.

The labor-market dictionary contains terms such as *employment*, *unemployment*, *labor market*, *job gains*, *payrolls*, *hiring*, *layoffs*, *wages*, *labor demand*, *labor supply*, *labor force*, *slack*, *participation*, *vacancies*, and *job openings*.

For a document d , the two attention scores are:

$$I(d) = 1000 \times \frac{\#\{\text{inflation terms in } d\}}{\#\{\text{words in } d\}}, \quad L(d) = 1000 \times \frac{\#\{\text{labor-market terms in } d\}}{\#\{\text{words in } d\}}. \quad (1)$$

The balance score is:

$$B(d) = I(d) - L(d). \quad (2)$$

A positive value means that inflation terms are more frequent than labor-market terms. Though, this measure also has limits: phrases can overlap with single words, and words such as *wages* can appear in both labor-market and inflation-pressure contexts.

Still, it is a simple proxy that allows us to get a general feeling of both subjects relative importance.

We then do simple checks because the dictionary is the core measurement choice.

The first check counts which terms drive the score. A dictionary measure can look stable while being driven by one repeated word, so the term composition matters.

The second check works at the sentence level. Each sentence is tagged as *inflation only*, *labor only*, *both*, or *neither*, depending on which dictionary terms it contains. This gives a sentence-level check on whether the document score is supported by identifiable sentences.

As a final document-level check, I use a low-rank TF-IDF representation. After vectorizing statements and minutes with unigrams and bigrams, I apply truncated singular value decomposition $X = U_k \Sigma_k V_k^\top$.

Intuitively, this decomposition looks for the main "directions" along which documents differ in their word usage.

The first components are the directions that explain the largest part of the variation in the TF-IDF matrix. If statements and minutes use systematically different vocabularies, lengths, formulas, or recurring expressions, this document-format effect should appear among the first components.

If the two document types separate clearly using this representation, then a large part of the textual variation comes from the format itself, before one even looks at the inflation-labor balance.

First diagnostics

Inflation language is more frequent than labor-market language in both document types. The difference is larger in statements. This is consistent with the role of the statement as a short policy message. Minutes are longer and cover more of the meeting, so the same macro topics take up a smaller share of the text.

Interestingly, we also observe an increase over the years in the use of labor related lexicon, for both types of documents.

The balance is usually positive, and the statement and minutes series move differently. Statements are often more inflation-dominant than minutes. The gap narrows in periods where labor-market language becomes more visible, especially around the first Covid shock. This is a clear examples where the labor side of the mandate appears more directly in the public statement.




Figure 2: Yearly average inflation and labor-market attention. Statements are more concentrated than minutes on both dimensions.

Table 2: Average mention scores per 1,000 words.

Document type	Inflation	Labor market	Balance
Statements	18.37	7.75	10.62
Minutes	10.58	5.98	4.60




Figure 3: Inflation-minus-labor balance by meeting. Faint lines show raw meeting scores; darker lines show a four-meeting rolling average.

Dictionary and sentence checks

Table 3: Most frequent dictionary terms in the corpus. Counts use the same matching rule as the score.

Document type	Topic	Most frequent terms
Statements	Inflation	inflation (1,340); inflation expectations (267); price stability (259); energy prices (76); core inflation (22)
Statements	Labor	labor market (410); employment (337); unemployment (138); job gains (64); slack (18)
Minutes	Inflation	inflation (10,387); inflation expectations (1,587); energy prices (985); price stability (779); core inflation (410)
Minutes	Labor	employment (2,293); labor market (2,235); unemployment (1,953); participation (444); slack (442)

The term check shows that the inflation side is strongly driven by the word *inflation*, which is expected. The next terms are also coherent expressions, which supports that the count is about FOMC language. On the labor side, the counts are spread across *labor market*, *employment*, and *unemployment*, with some use of *job gains*, *participation*, and *slack*.




Figure 4: Sentence-level dictionary check. Each sentence is classified as inflation only, labor only, both, or neither.

The sentence check supports the generic expectation: statements more often compress both sides of the mandate into a short text, while minutes spread the discussion across many more sentences.

Decisions, periods, and document gaps

The balance score is related to rate decisions in the expected direction. Hike meetings have higher average balances than cuts or holds. This descriptive check says that meetings near hikes use more inflation language relative to labor-market language. To avoid any confusion, the points shown are not independent at all and there is significant correlation from one meeting to the following, leading to clusters in the visualization.




Figure 5: Mandate balance by current and next rate decision. Points are meetings; diamonds are group means.

The period comparison is coarse, but it helps locate the main episodes. The Covid shock has a lower balance, mainly because labor-market language becomes more visible in statements. The global financial crisis remains more positive in this dictionary, probably because much of that discussion uses financial-stability, credit, and risk language outside the labor terms listed above.




Figure 6: Average mandate balance by period. Sample sizes are shown under the period labels.

These events highlight how meeting-specific the score can be. A large negative gap means that the statement was much more one-sided than the minutes, while the two documents can still tell a consistent story about the same meeting.




Figure 7: Minutes minus statement balance. Negative values mean that statements are more inflation-dominant than minutes.

0.1 Singular Value Decomposition projection




Figure 8: TF-IDF/LSA map of statements and minutes. Color is the inflation-labor balance.

The first two SVD dimensions separate statements and minutes very clearly. In this low-rank TF-IDF map, document format is a major source of variation. The two components even separate the statements in two very distinct clusters, with one having especially higher inflation-labor balance points (in red).

From this graph, it seems some sub-category of statements are disproportionately inflation heavy.

Sentence-Level Weak Labels and Topic Classification

The preliminary exploration was the simplest, first approach analysis to get a feel for the data, and relied mostly on dictionaries.

However, FED communication is way more subtle and calculated. We would like to have a finer view of what's being communicated and how it correlates to its' mandates. Indeed, a sentence about credit

conditions, market liquidity, or financial stress can be central to the Fed’s communication while using none of the inflation or labor-market words above.

The global financial crisis is the clearest example: much of the relevant language is about markets, credit, and liquidity, which sit outside the two mandate dictionaries.

This part should be read as a semantic audit rather than as a new ground truth. The LLM-assisted labels are useful because they impose a consistent topic schema on a small sample of sentences, including sentences missed by the dictionary. They are still weak labels. The models below are therefore evaluated by their agreement with those labels, not by their accuracy against human-coded truth.

To check this more directly, I split the statements and minutes into a total of 58,976 sentences. For each sentence, I keep the meeting date, document type, sentence position, rate decision, period, election-window indicators, and dictionary label.

I then draw a balanced sample of 3000 sentences. The sample includes statements and minutes, inflation-only sentences, labor-only sentences, sentences containing both dictionaries, sentences containing neither dictionary, meetings with large statement–minutes gaps, crisis and post-Covid inflation periods, and federal election windows.

Automatic sentence labels

The 3000 sampled sentences are labelled with a local language model (llama3.2) using a fixed JSON schema. The main topic can be *inflation*, *labor*, *both*, *financial conditions*, *policy decision*, or *other*. The prompt asks the model to base the classification on the sentence itself and to use *financial conditions* for sentences about credit, markets, liquidity, volatility, asset prices, or market functioning. These labels are pseudo-labels: they provide a second, broader reading of the same sentences and a training target for a small classifier.

In this balanced sample, the most striking result is how often the automatic label is *financial conditions*. It reveals an important category of sentences outside the inflation-labor dictionary. In the pseudo-labels, many dictionary-*neither* sentences contain substantive background on credit, market functioning, or financial stress.

Dictionary labels versus weak LLM sentence labels

Dictionary sentence label						
	inflation	labor	both	financial_conditions	policy_decision	other
	inflation	labor	both	financial_conditions	policy_decision	other
	inflation	labor	both	financial_conditions	policy_decision	other
inflation	250	1	22	131	1	7
labor	1	137	91	41	1	20
both	30	1	103	62	1	9
other	14	19	48	1116	163	731
LLM main topic						

Figure 9: Dictionary sentence labels compared with pseudo-labelled sentence topics. The main gap is that many dictionary-*other* sentences are labelled as financial conditions.

The comparison is reassuring in the easy cases. Sentences with explicit inflation words are mostly labelled as inflation or both. Sentences with explicit labor-market words are mostly labelled as labor or both. Financial conditions appear across all dictionary buckets, especially in sentences that the dictionary calls *other*.




Figure 10: Embedding classifier on the pseudo-labelled sentence sample no-training baseline: each sentence is assigned to the nearest topic prototype.

I then train small sentence classifiers on the pseudo-labels. The question is direct: can the sentence topics be learned from text, and do embeddings add information beyond exact dictionary matching? The target is the main topic, with *other* and *policy decision* merged into a single non-mandate category when the class sizes are too small.

I compare three models. The first is the dictionary rule itself. It is transparent and uses sentences that contain inflation or labor terms. The second is a TF-IDF logistic-regression classifier with unigrams and bigrams. This tests whether a standard sparse text representation can learn the pseudo-labels. The third is a logistic-regression classifier on local sentence embeddings from `nomic-embed-text`. This tests whether semantic features capture nearby language, for example market stress or liquidity language, outside the dictionary. I also include a no-training embedding baseline: each sentence embedding is compared with short prototype descriptions of inflation, labor, financial conditions, policy decisions, and other text.




Figure 11: Sentence-level classifier scores. Macro-F1 is the main metric because the pseudo-labels are unevenly distributed.




Figure 12: Confusion matrix for the embedding logistic-regression sentence classifier.

Table 4: Sentence-topic classification against LLM-assisted weak labels.

Model	Accuracy	Macro-F1
Dictionary rule	0.535	0.526
TF-IDF + logistic regression	0.718	0.724
Embedding + logistic regression	0.768	0.765

The comparison gives the expected ordering. The dictionary rule is useful but limited by exact matching. The TF-IDF classifier improves substantially, and the embedding classifier performs slightly better again. These numbers should not be read as accuracy against a human-coded ground truth. They measure agreement with the LLM-assisted weak labels. The useful result is comparative: sentence topics contain recoverable structure beyond exact dictionary hits, and embeddings help capture some of that structure.

0.2 Applying the classifier to the corpus

After fitting the selected sentence classifier, I apply it to the full sentence corpus. When full sentence embeddings are available, I use the embedding classifier. Otherwise, the notebook falls back to the TF-IDF classifier. In the results reported here, full embeddings are available and the full corpus is classified with the embedding logistic model. The resulting file `data/processed/fomc_sentence_topic_predictions.csv` keeps the predicted topic, prediction confidence, date, year, document type, period, and election-window variables for each sentence.




Figure 13: Predicted sentence-topic shares for statements and minutes. The classifier view keeps the inflation–labor contrast but adds financial conditions and policy-decision language.




Figure 14: Predicted sentence-topic shares over time. The figure focuses on inflation, labor, and financial conditions to keep the plot readable.




Figure 15: Dictionary inflation-minus-labor balance against the classifier inflation-minus-labor sentence share. Each point is a meeting-document pair.

This full-corpus step is the bridge between the old score and the new classifier. The dictionary balance remains important because it has a direct interpretation. The classifier adds a second view: whether the same meeting looks inflation-heavy or labor-heavy after sentences are classified by broader semantic features. In the full corpus, the classifier predicts a large financial-conditions share, especially in minutes: 49.4% of minutes sentences against 24.6% of statement sentences. Statements have more sentences that combine both sides of the mandate, 17.9% against 5.0% in minutes.

The absolute levels should be read cautiously because the classifier is trained on a small, deliberately balanced weak-label sample. The more useful comparison is relative: the model assigns a larger financial-conditions share to minutes than to statements, and it assigns a larger combined-mandate share to statements than to minutes. This is consistent with the document-format interpretation: statements compress the public message, while minutes contain more background discussion.

0.3 Election and institutional timing

Finally, I add a small calendar check. I mark presidential election years, midterm years, meetings within 90 days of a federal election, the first meeting of each year, even-year Board-cycle years, and Chair appointment or reappointment years. These variables are descriptive calendar markers for comparing topic mix around scheduled institutional and electoral dates.




Figure 16: Dictionary balance inside and outside 90-day federal election windows. Points are meetings; diamonds are group means.




Figure 17: Predicted sentence-topic shares inside and outside 90-day federal election windows.

The reason to include this check is simple. FOMC communication is public and scheduled, and some meetings occur close to major political events or institutional calendar points. A visible shift

in the topic mix around those dates would be worth noticing. Here the comparison is descriptive, with a small sample and coarse timing. In this sample, the differences are small: the predicted topic shares inside and outside 90-day election windows are almost identical, and the dictionary balance changes little. This is not evidence about political influence. It is a negative descriptive check around scheduled calendar windows.

1 Discussion

The main result is simple. Statements and minutes allocate attention differently. Statements are short and more inflation-dominant. Minutes are longer, more diffuse, and closer to balance. The statement–minutes gap is itself informative because it shows how much the immediate policy message differs from the longer meeting account.

The point is not established by one number only. The dictionary score gives the most transparent measure. The term audit shows which expressions drive the score. The sentence audit checks whether document-level balances correspond to identifiable sentences. The low-rank TF–IDF map shows that document format is a major source of variation. The weak-label sentence classifier then adds a semantic check: broader topics, especially financial conditions, can be recovered from sentence text even when exact inflation or labor-market terms are absent. These checks do not prove a causal story, but they make the descriptive result more credible.

The result is also expected. Inflation is the part of the dual mandate most directly tied to the Fed’s credibility and to expectations. That makes it natural for statements to put inflation language in the foreground, especially when the Committee is explaining the policy decision. Labor-market language is still present, but it often works differently: it helps describe slack, demand, wages, and the cost of tightening. The two sides of the mandate are communicated through different textual roles.

The sentence-level results show why the dictionary should not be the whole analysis. In some periods, especially stress periods, many important sentences are about financial conditions, credit, market functioning, and liquidity. These sentences are relevant to FOMC communication, but they are not inflation or labor-market sentences in the narrow dictionary sense. The balance score is therefore best read as a descriptive measure of inflation-versus-labor attention, while the sentence classifier adds a third dimension that explains part of what the dictionary leaves out.

The rate-decision and period figures move in plausible directions. Hike meetings are more inflation-heavy than cut meetings, and the Covid shock is one of the clearest points where labor-market language becomes more visible. These are associations in a small sample, not causal estimates. The election-window check is similar. It is included as a calendar comparison, and the result is mostly negative: in this sample, the topic mix inside and outside 90-day federal election windows looks very similar.

The main methodological lesson is that document format matters. Models that pool statements and minutes can mainly learn the difference between a short public statement and a long meeting summary. This is why the report keeps the dictionary score separated by document type and uses sentence models as a check rather than as a replacement. The methods point in the same general direction, but each has a different weakness. The conclusion should therefore be read as a triangulated descriptive finding, not as a single definitive measurement.

2 Connection with the Companion Report

This report and the companion hawkish/dovish report share the same corpus, but they measure different objects. The mandate balance asks what macro topic is emphasized: inflation or the labor market. The hawkish/dovish score asks what policy signal the text sends. Keeping the two dimensions separate avoids treating all inflation language as hawkish and all labor-market language as dovish.

3 Conclusion

In this sample of 199 FOMC meetings, statements and minutes give different textual weights to the two sides of the dual mandate. Statements are much shorter and more inflation-dominant. Minutes are more balanced and contain a wider discussion. This is the main descriptive result of the report.

Several checks point in the same direction. The dictionary balance is higher in statements. The term audit shows that the result is driven by economically meaningful expressions, even though the word *inflation* plays a large role. The sentence-level dictionary audit shows that statements more often compress inflation and labor-market language into the same short document. The low-rank TF-IDF map shows that statements and minutes are different textual formats. The weak-label sentence classifier adds one more check: broader sentence topics are partly learnable from text, and embedding features agree more closely with the weak labels than the dictionary rule.

The qualification is important. The inflation–labor contrast does not cover all of FOMC communication. In stress periods, financial conditions become central, and a pure dual-mandate dictionary misses that part of the story. The classifier view helps make this visible, especially in minutes. Its absolute topic shares should be read cautiously because the model is trained on a small weakly labelled sample, but the comparison with the dictionary is useful.

The report therefore supports a cautious conclusion. Public FOMC communication in this sample puts more relative inflation emphasis in statements than in minutes. Minutes give a broader account of the meeting, including more labor-market and financial-conditions language. This does not identify the Fed’s objective function, and it does not make a causal claim about policy or politics. It is a descriptive NLP result: different public documents from the same meetings allocate textual attention differently, and simple machine-learning tools help check where the transparent dictionary measure is reliable and where it is incomplete.

A TF-IDF Sanity Check




Figure 18: Macro-F1 scores for TF-IDF models. The dashed line is the majority-class baseline.

A small TF-IDF classification exercise remains in the appendix. The target is either the current rate decision or the next meeting’s rate decision. The models are standard TF-IDF+logistic regression and TF-IDF+linear SVM classifiers, compared with a majority-class baseline. The split is chronological: the first 70% of meetings are used for training and the last 30% for testing.

The rate-decision scores remain modest. The majority baseline has macro-F1 0.222. For the current decision, the best model is TF-IDF+linear SVM on full minutes, with macro-F1 0.418. For the next decision, the best model is again based on minutes, with macro-F1 0.396.

This exercise is a small appendix check. The sample has 199 meetings, the classes are unbalanced, and minutes are released after the meeting. The main takeaway is smaller: minutes contain more decision-related context than statements, even though statements have a higher density of mandate terms.

B Document-Level Embedding Check




Figure 19: Dictionary balance against embedding prototype balance.




Figure 20: PCA map of local document embeddings. Color is the embedding prototype balance.

Before the sentence-level model, I also tried a small local embedding check with `nomic-embed-text`. Each document is represented by a short beginning/middle/end excerpt, and the document vector is compared with one inflation prototype phrase and one labor-market prototype phrase. This document-level view is noisier than the sentence classifier, especially for minutes, so I keep it in the appendix. It points to the same message as the LSA map: statements and minutes are different document formats. The main embedding result in the report is now the sentence-level classifier in section 0.1.

References

- Acosta, M., and Meade, E. E. Hanging on every word: Semantic analysis of the FOMC's postmeeting statement. *FEDS Notes*, Board of Governors of the Federal Reserve System, 2015.
- Blinder, A. S., Ehrmann, M., Fratzscher, M., De Haan, J., and Jansen, D.-J. Central bank communication and monetary policy: A survey of theory and evidence. *Journal of Economic Literature*, 46(4):910–945, 2008.
- Deerwester, S., Dumais, S. T., Furnas, G. W., Landauer, T. K., and Harshman, R. Indexing by latent semantic analysis. *Journal of the American Society for Information Science*, 41(6):391–407, 1990.
- Doh, T., Kim, S., and Yang, S.-K. How you say it matters: Text analysis of FOMC statements using natural language processing. *Economic Review, Federal Reserve Bank of Kansas City*, 106(1):25–40, 2021.
- Doh, T., Song, D., and Yang, S.-K. Deciphering Federal Reserve communication via text analysis of alternative FOMC statements. *American Economic Journal: Macroeconomics*, forthcoming.
- Gentzkow, M., Kelly, B., and Taddy, M. Text as data. *Journal of Economic Literature*, 57(3):535–574, 2019.
- Hansen, S., McMahon, M., and Prat, A. Transparency and deliberation within the FOMC: A computational linguistics approach. *The Quarterly Journal of Economics*, 133(2):801–870, 2018.
- Loughran, T., and McDonald, B. When is a liability not a liability? Textual analysis, dictionaries, and 10-Ks. *Journal of Finance*, 66(1):35–65, 2011.
- Shapiro, A. H., and Wilson, D. J. Taking the Fed at its word: A new approach to estimating central bank objectives using text analysis. *The Review of Economic Studies*, 89(5):2768–2805, 2022.
- Tadle, R. C. FOMC minutes sentiments and their impact on financial markets. *Journal of Economics and Business*, 118:106021, 2022.